

Neural Network Models for Object Recognition

Group 1: Presentation



Contents

Neural Network Design

Aim

Data split

Preprocessing - Encoding to category for the 10 classes

Model design

Overfitting alarm

Dropout technique

Model evaluation - kernel size

Model evaluation - No of kernels & Early Stopping

Epochs used

Model evaluation- no.of convolution layers

Discussion

Conclusion

References

Code



Introduction

Object recognition - deep learning application

Machines perceive like humans - drive cars, read and comprehend documents

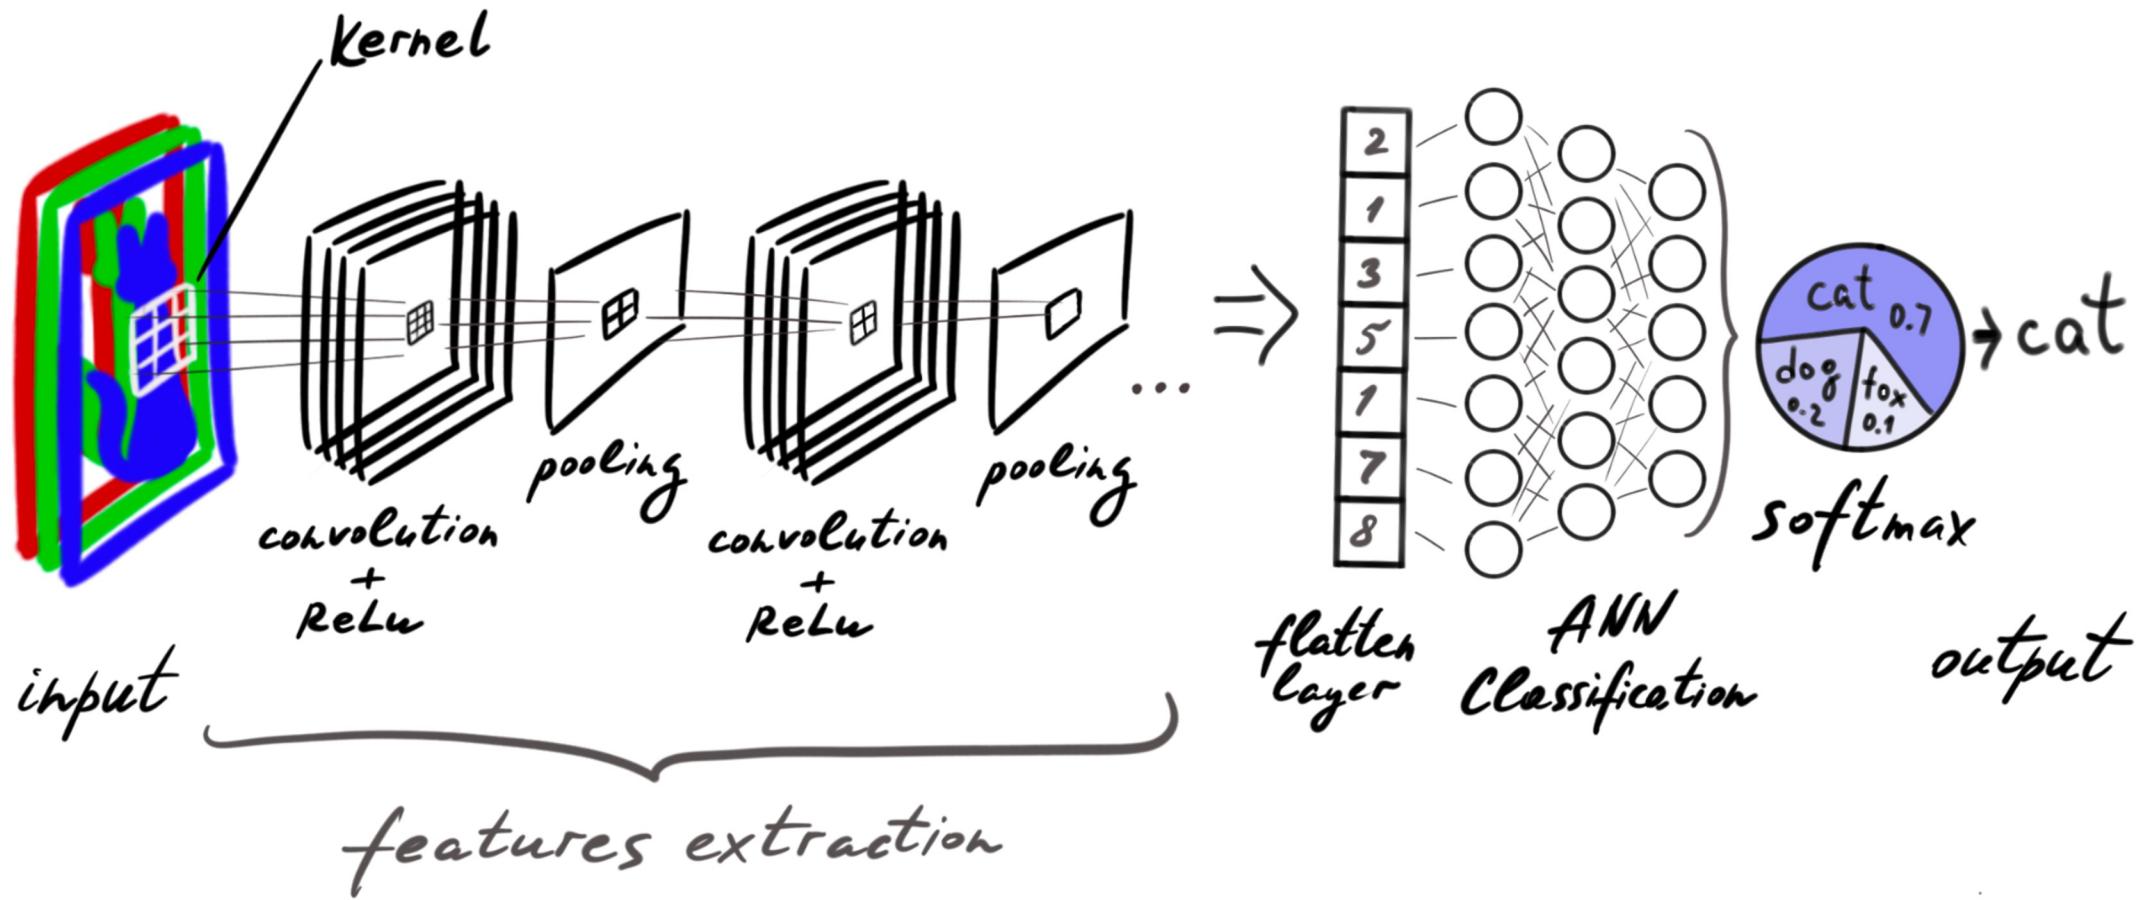
Convolution neural network (CNN) - object detection through feature learning

Artificial neural network (ANN) - Classify object



Neural Network Design

Two main steps: Feature learning and classification



Aim

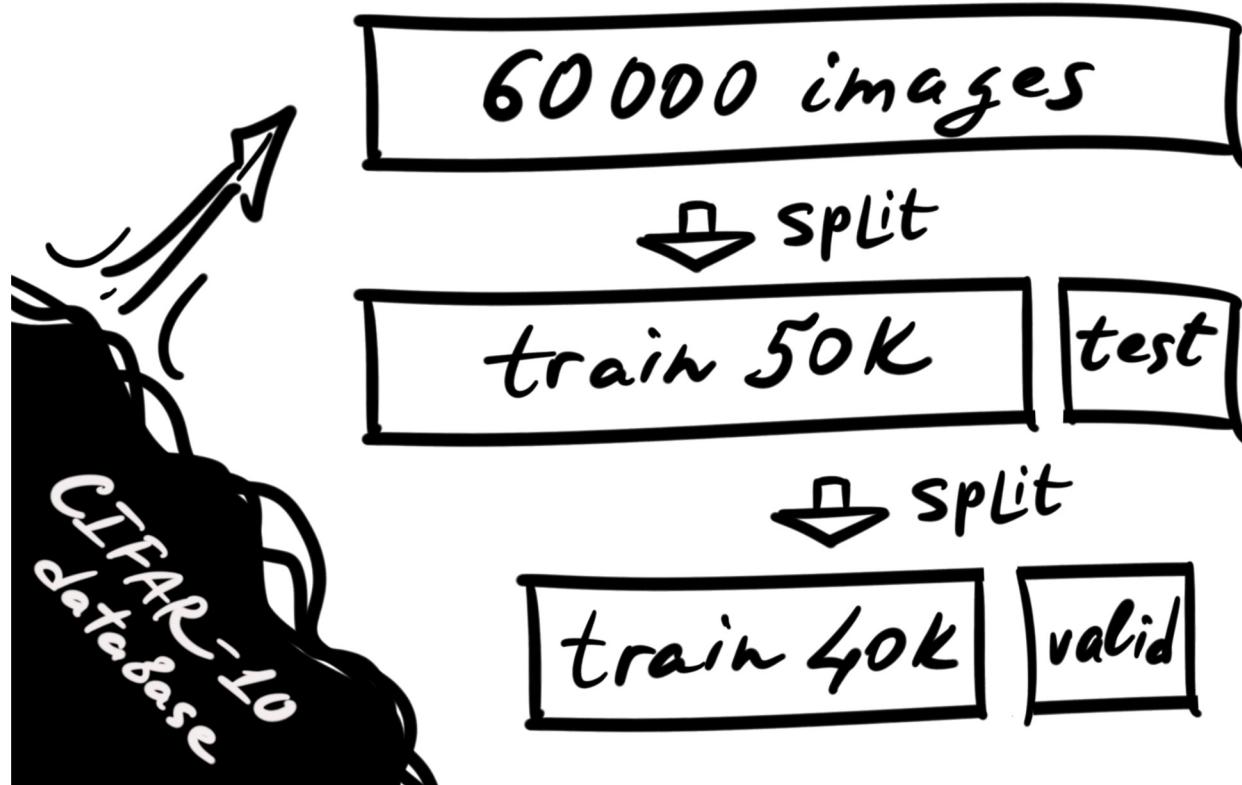
To build a neural network for object recognition

To evaluate the built model

Image data: Open source CIFAR-10- Object Recognition image dataset, hosted in Kaggle. It consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. 10 classes: airplane, automobile, bird, cat, dear, dog, frog, horse, ship, truck

Python libraries used: numpy, tensorflow, keras, pandas, seaborn,

Data split



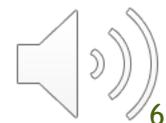
The CIFAR-10 dataset has 50000 training images and 10000 test images. 10000 images out of 50000 training images were used for validation.

Training:Validation:Testing = 40000:10000:10000

Training data: Model trained the data

Validation data: Model fine-tuning, help to improve model performance

Testing data: Provides measure of accuracy of built model



Preprocessing - Normalising the data

```
In [26]: x_train = x_train / 255.0
In [27]: x_test = x_test / 255.0

In [21]: x_train[0]
Out[21]: array([[[ 59,  62,  63],
   [ 43,  46,  45],
   [ 50,  48,  43],
   ...,
   [158, 132, 108],
   [152, 125, 102],
   [148, 124, 103]],

  [[ 16,  20,  20],
   [  0,   0,   0],
   [ 18,   8,   0],
   ...,
   [123,  88,  55],
   [119,  83,  50],
   [122,  87,  57]],

  [[ 25,  24,  21],
   [ 16,   7,   0],
   [ 49,  27,   8],
   ...,
   [118,  84,  50],
   [120,  84,  50],
   [109,  73,  42]],

  ...,
```



```
In [58]: x_train
Out[58]: array([[[[0.1372549 , 0.09803922, 0.10196078],
   [0.10588235, 0.08235294, 0.08235294],
   [0.09803922, 0.07843137, 0.0745098 ],
   ...,
   [0.51764706, 0.50588235, 0.50588235],
   [0.52156863, 0.4745098 , 0.45490196],
   [0.49411765, 0.45098039, 0.44313725]],

  [[[0.24705882, 0.21568627, 0.19607843],
   [0.1254902 , 0.10588235, 0.08235294],
   [0.06666667, 0.05098039, 0.03137255],
   ...,
   [0.4       , 0.37254902, 0.34509804],
   [0.41176471, 0.34901961, 0.29803922],
   [0.39215686, 0.3372549 , 0.30196078]],

  [[[0.38823529, 0.35686275, 0.32941176],
   [0.19215686, 0.17647059, 0.14509804],
   [0.05882353, 0.04705882, 0.01960784],
```

With Numpy library providing normalisation for every pixel in both training and testing datasets.

From 0 to 255 for each of three channels (RGB) it comes to a range between 0 and 1.



Preprocessing - Encoding to_category for the 10 classes

```
y_cat_train_all = to_categorical(y_train_all, 10)
```

```
array([[3],  
       [8],  
       [8],  
       ...,  
       [5],  
       [1],  
       [7]], dtype=uint8)
```

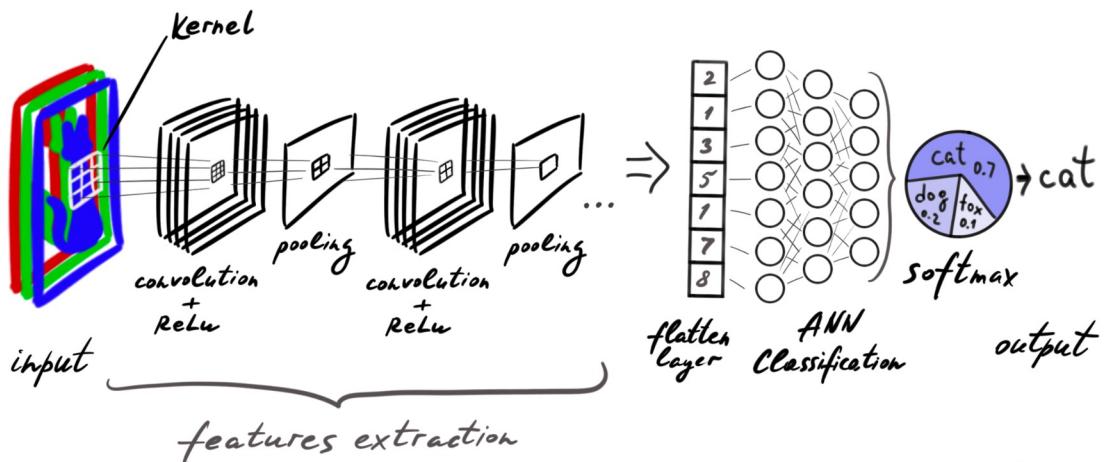


```
array([[0., 0., 0., ..., 0., 0., 0.],  
       [0., 0., 0., ..., 0., 0., 1.],  
       [0., 0., 0., ..., 0., 0., 1.],  
       ...,  
       [0., 0., 0., ..., 0., 0., 1.],  
       [0., 1., 0., ..., 0., 0., 0.],  
       [0., 1., 0., ..., 0., 0., 0.]], dtype=float32)
```



Model design

Schematic image for the neural network:



The same written in code:

```
model = Sequential()

# First CONVOLUTIONAL LAYER
model.add(Conv2D(filters=32, kernel_size=(4,4),input_shape=(32, 32, 3), activation='relu',))
# POOLING LAYER
model.add(MaxPool2D(pool_size=(2, 2)))

# Second CONVOLUTIONAL LAYER
model.add(Conv2D(filters=32, kernel_size=(4,4),activation='relu',))
# POOLING LAYER
model.add(MaxPool2D(pool_size=(2, 2)))

# FLATTEN IMAGES
model.add(Flatten())

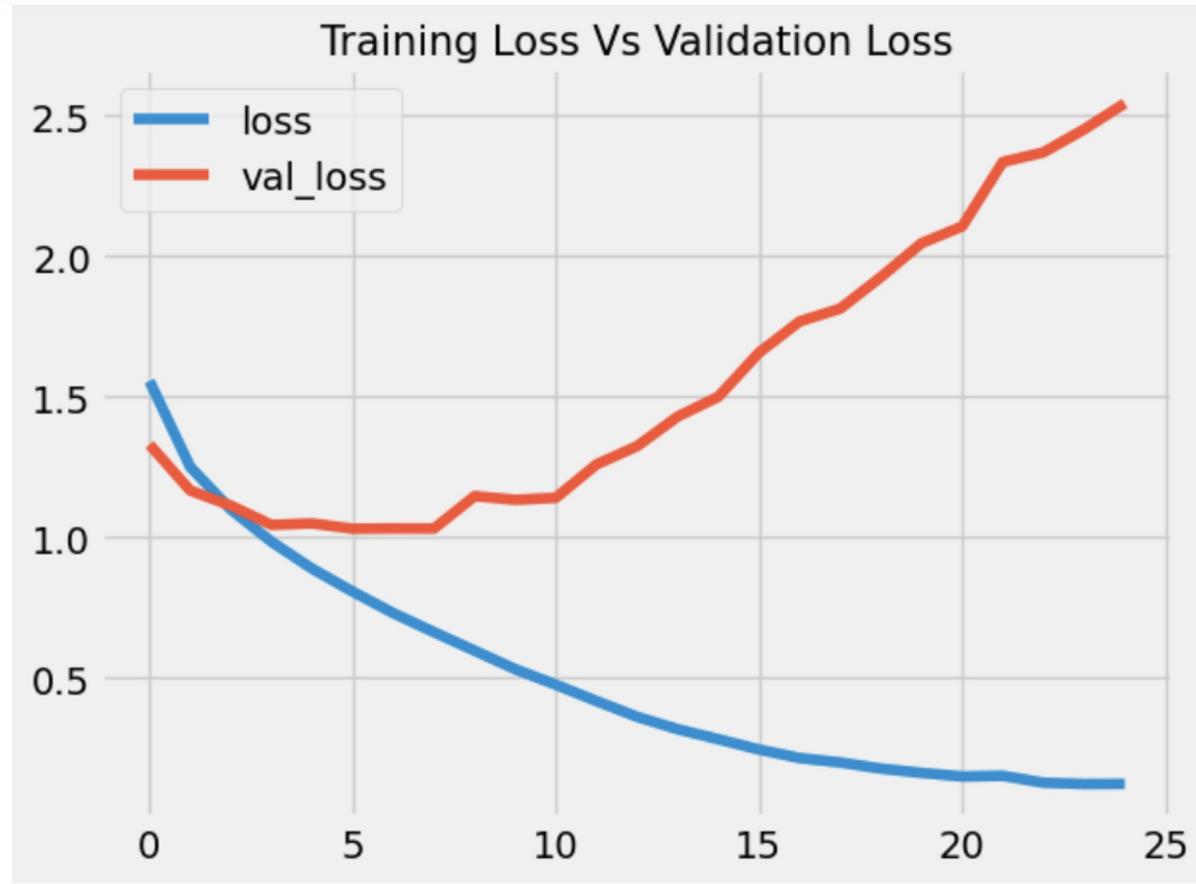
# Condense
model.add(Dense(256, activation='relu'))

# classifier
model.add(Dense(10, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```



Overfitting alarm



It starts well, but after a few first iterations the loss in validation set increases.

The model shows an overfitting relationship with a very low training loss but relatively high validation loss.

Overfitted neural networks are not good at generalizing. (Doon et al., 2018)



Dropout technique

*Without
Dropout*



98%

Dropout 0.25



86%

Dropout 0.5



74%

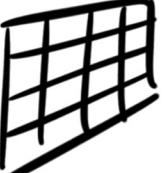
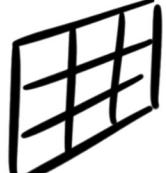
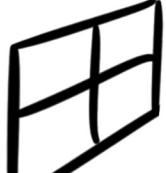
training accuracy

Using the Dropout technique to manage with overfitting.
It randomly drops units from the network with some probability during training. (Srivastava et al., 2014)

However on the long distance of many epochs Dropout decreases the training accuracy.



Model evaluation - kernel size

Kernels:	Validation accuracy:	Training accuracy:
4×4 	67%	95%
3×3 	69%	97%
2×2 	70%	98%

A little correlation of a kernel size and validation/training accuracy of the model.

A relatively small size of the kernel helps to find more details in the picture. The bigger kernel size could generalise better.

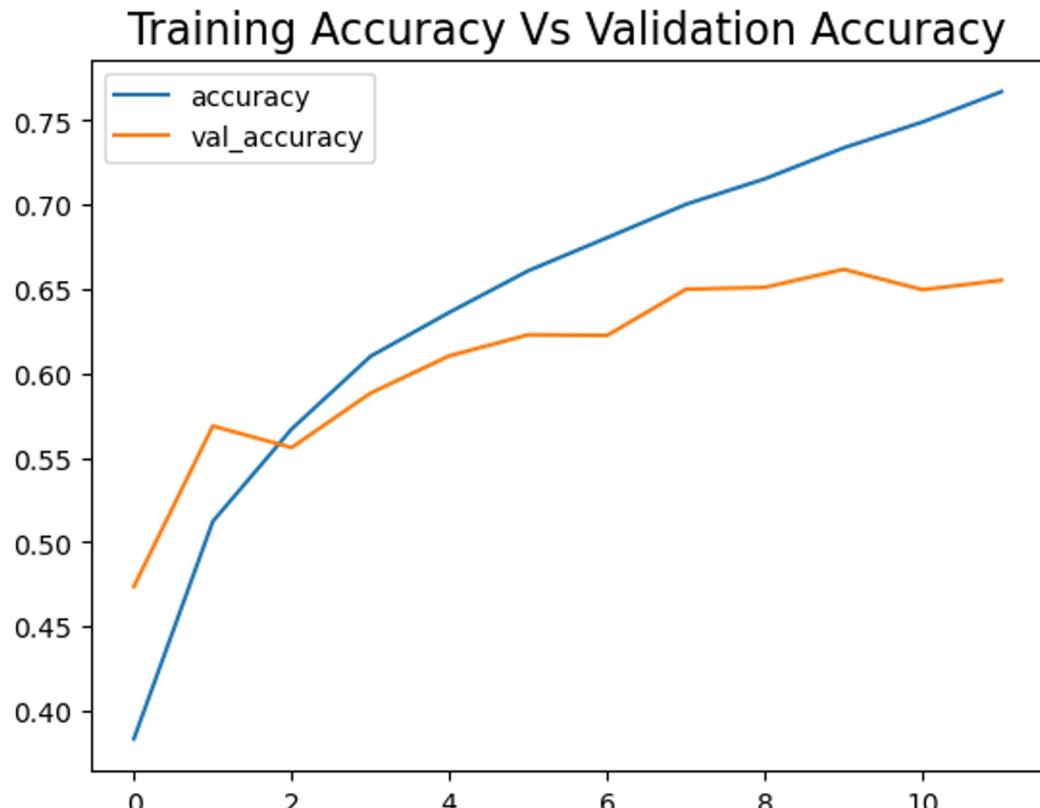
(Cortes et al., 2010)



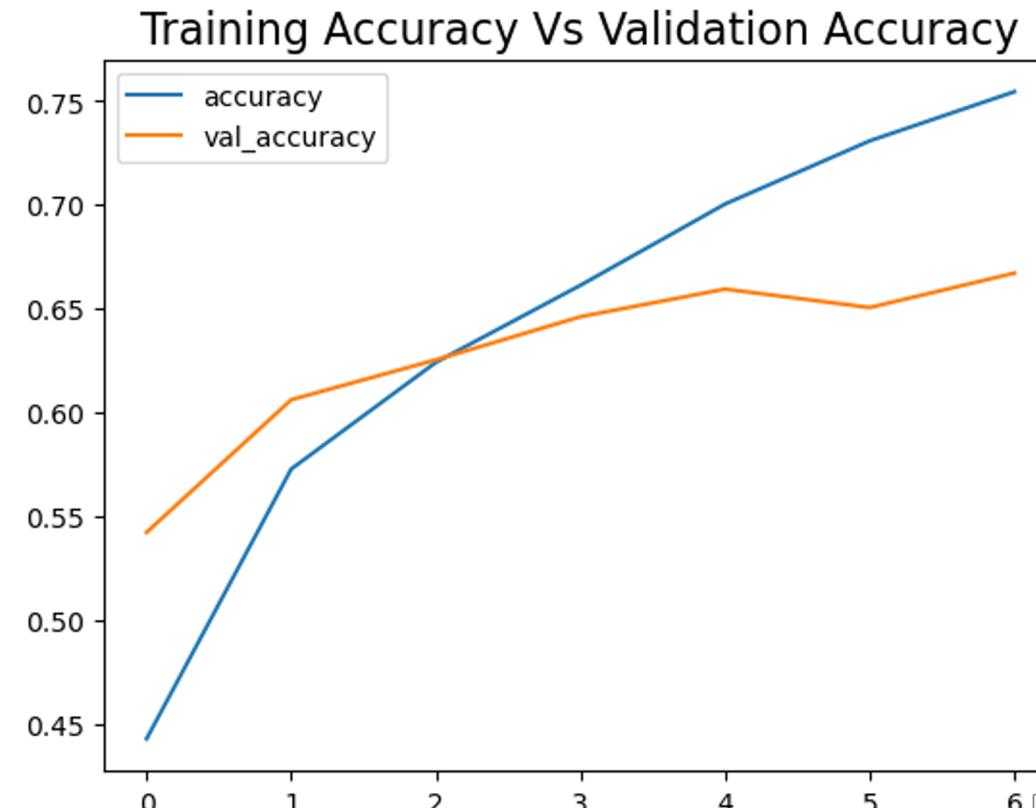
Model evaluation - No of kernels & Early Stopping

Kernels 32 - Training accuracy 75.5%

Kernels 64 - Training accuracy 76.7%



Another way to avoid overfitting



Epochs used

Our model used 12 epochs - optimal

Used early stopping, patience 2

Too many epochs - overfitting problem

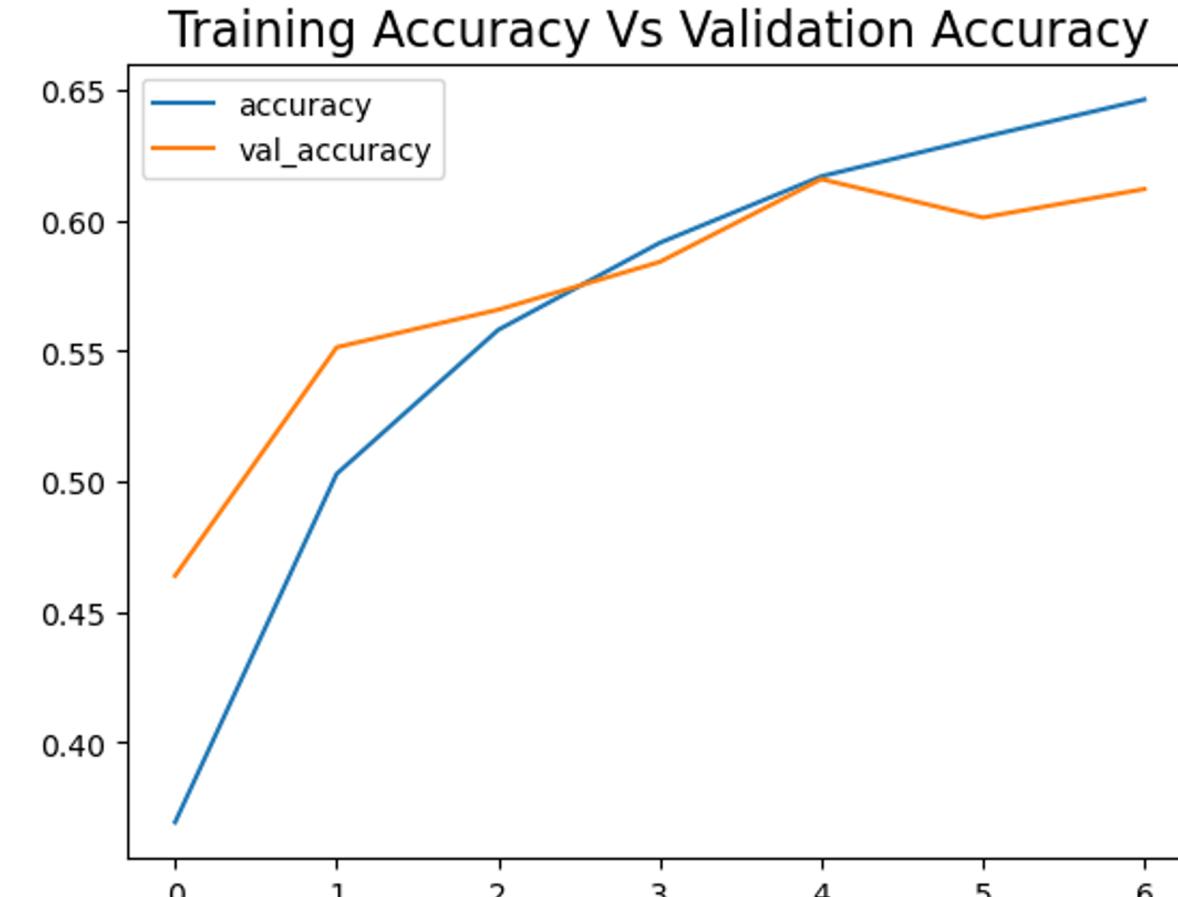
	loss	accuracy	val_loss	val_accuracy
0	1.680713	0.383500	1.433848	0.4738
1	1.364287	0.512325	1.236065	0.5689
2	1.221159	0.567125	1.268536	0.5561
3	1.111533	0.610225	1.172620	0.5884
4	1.031733	0.636300	1.119634	0.6105
5	0.966625	0.660950	1.091673	0.6230
6	0.904833	0.680550	1.109121	0.6226
7	0.850928	0.700350	1.017413	0.6500
8	0.802035	0.715475	1.037993	0.6511
9	0.755995	0.733825	1.014872	0.6618
10	0.711462	0.749150	1.078260	0.6497
11	0.662567	0.767125	1.082580	0.6553



Model evaluation- no.of convolution layers

Adding one more convolution layer 2+1

Validation accuracy more closer to training accuracy



Discussion

Key findings:

- The smaller kernel size increases the accuracy for the training model
- Double number of kernels provides better accuracy (+1.2%)
- Dropout helps to manage overfitting, but decreases further accuracy
- One more *Convolutional + ReLU* layer increases accuracy
- Early stopping trick prevents from overfitting on the long distance



Conclusion

```
#predictions = model.predict_classes(x_test)
predictions = np.argmax(model.predict(x_test), axis=-1)

313/313 [=====] - 1s 3ms/step

print(classification_report(y_test,predictions))

precision    recall    f1-score   support
0            0.72      0.77      0.74      1000
1            0.81      0.81      0.81      1000
2            0.65      0.60      0.62      1000
3            0.50      0.61      0.55      1000
4            0.67      0.66      0.67      1000
5            0.65      0.52      0.58      1000
6            0.81      0.75      0.78      1000
7            0.75      0.79      0.77      1000
8            0.81      0.82      0.81      1000
9            0.79      0.77      0.78      1000

accuracy                           0.71      10000
macro avg       0.71      0.71      0.71      10000
weighted avg    0.71      0.71      0.71      10000
```

The CIFAR-10 data set is divided into a training and a test set, but does not come with a validation set by default.

We created a validation set, then we trained the model on the training set, and tested on the validation set to see if it is a good fit.

In the end the testing set was used to calculate accuracy: **about 71%**.



References

- The CIFAR-10 dataset. Available at: <https://www.cs.toronto.edu/~kriz/cifar.html> [Accessed 1 Dec 2022]
- Krizhevsky, A. & Hinton, G. (2009) Learning multiple layers of features from tiny images.
- Recht, B., Roelofs, R., Schmidt, L. and Shankar, V. (2018) Do cifar-10 classifiers generalize to cifar-10?. arXiv preprint arXiv:1806.00451.
- Doon, R., Rawat, T.K. and Gautam, S. (2018). Cifar-10 classification using deep convolutional neural network. In 2018 IEEE Punecon (pp. 1-5). IEEE.
- Krizhevsky, A. and Hinton, G. (2010) Convolutional deep belief networks on cifar-10. Unpublished manuscript, 40(7), pp.1-9.
- Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I. and Salakhutdinov, R. (2014) Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research*, 15(1), pp.1929-1958.
- Cortes, C., Mohri, M. and Talwalkar, A. (2010) On the impact of kernel approximation on learning accuracy. In Proceedings of the thirteenth international conference on artificial intelligence and statistics (pp. 113-120). JMLR Workshop and Conference Proceedings.

Code

- Please find the code added at the end of the transcript